

Capstone Lab: Full Implementation of a Graph-Powered GenAI Application

◆ Step 1: Define the Problem & Use Case

Objective: Choose a real-world problem where semantic understanding, retrieval, and reasoning matter.

Example Use Cases:

- Healthcare: Clinical decision support
- Retail: Personalized shopping assistant
- Legal: Contract clause summarization

Deliverables:

- Problem statement
 - Business objective (what value GenAI brings)
 - Key actors and workflows (input/output)
-

◆ Step 2: Design Ontology & Build Knowledge Graph

Tools: Protégé, TopBraid Composer, RDF/OWL, Neo4j/Stardog

Actions:

1. Define domain-specific ontology (classes, relationships, properties)
2. Save in RDF/OWL format
3. Populate the graph with sample or real data
4. Import RDF into Neo4j/Stardog using a loader or converter

Deliverables:

- Ontology diagram (PDF or image)
 - Graph snapshot (nodes, edges)
 - Example SPARQL/Cypher queries
-

◆ Step 3: Build RAG Pipeline with Graph + Vector Retrieval

Tools: FAISS or Pinecone (for vector search), RDF graph database (Neo4j), OpenAI API

Actions:

1. Embed your knowledge snippets (documents, FAQs) using OpenAI or HuggingFace embeddings
2. Store embeddings in FAISS or Pinecone
3. Build SPARQL/Cypher queries to retrieve entities from your graph
4. Merge both results into a unified context for the prompt

Deliverables:

- Hybrid retriever function (code)
 - Example retrieval results (log file or screenshot)
 - Combined context passed to LLM
-

◆ **Step 4: Implement Agent Workflow (LangGraph or CrewAI)**

Agent Roles:

- **Planner Agent:** Breaks down user input into subtasks
- **Retriever Agent:** Queries graph and vector DBs
- **Summarizer Agent:** Synthesizes response from retrieved content

Actions:

1. Use LangGraph or CrewAI to create each agent node
2. Define memory (context passing), tools (SPARQL, FAISS client), and transitions
3. Chain agents into a LangGraph workflow

Deliverables:

- Agent definitions (JSON or Python)
 - Workflow diagram (Mermaid or [Draw.io](#))
 - Demo output from a sample query
-

◆ **Step 5: Containerize and Deploy to the Cloud**

Platforms: AWS Fargate, Azure Container Apps, or GCP Cloud Run

Actions:

1. Write Dockerfile for your app (LLM + graph + vector + agent orchestrator)
2. Push image to ECR/ACR/GCR
3. Deploy with proper secrets for APIs and DBs
4. Set autoscaling, logging (CloudWatch or Prometheus), and endpoint routing

Deliverables:

- Dockerfile + cloud deploy config (YAML/JSON)
 - Live endpoint or screencast of deployed app
 - Monitoring dashboard screenshot
-

◆ **Step 6: Document the Architecture**

Tools: Mermaid, [Draw.io](#), Markdown, Google Docs

Sections to Include:

- Business problem and user need
- Ontology structure and graph schema
- RAG pipeline with code snippets
- Agent flow with reasoning logic
- Deployment diagram and endpoints

Deliverables:

- Capstone documentation PDF or GitHub repo
 - Explainer video or slides (optional)
-

 **Step 7: Stakeholder Presentation**

Actions:

1. Prepare a 10-15 minute walkthrough
2. Cover what problem you solved, how it works, and what value it provides
3. Show live demo or screencast
4. Invite questions, feedback, and reflection

Deliverables:

- Slide deck (PPT or PDF)
 - Recording or live presentation
 - Reviewer feedback summary
-

 **Final Output**

- Working, deployed GenAI application with hybrid retrieval and agents
- Knowledge graph and ontology implementation
- Complete documentation and demo materials

Congratulations—you've built a production-grade, graph-powered Generative AI system that integrates symbolic reasoning, vector search, and cloud-native deployment.